

Coordination and Agreement

➤ Introduction

- *How processes coordinate their actions?*
- Main assumptions in coordination:
 - ✓ Each pair of processes is connected by reliable channels
 - ✓ Processes independent from each other
 - ✓ Processes fail only by crashing

Coordination and Agreement

➤ Distributed Mutual Exclusion

- Distributed processes require a mechanism that can coordinate their activities because they share a resource or collection of resources
- Mutual exclusion is required to
 - ✓ prevent interference
 - ✓ ensure consistency when accessing the resources



10 11 12 13

Coordination and Agreement

➤ Distributed Mutual Exclusion

- Distributed processes require a mechanism that can coordinate their activities because they share a resource or collection of resources
- Mutual exclusion is required to
 - ✓ prevent interference
 - ✓ ensure consistency when accessing the resources



Coordination and Agreement

➤ Distributed Mutual Exclusion

■ Algorithms for mutual exclusion

✓ Requirements for mutual exclusion are:

- Safety - At most one process may execute in the critical section (CS) at a time.
- Liveness - Requests to enter and exit the critical section eventually succeed.
- Ordering - If one request to enter the CS happened-before another, then entry to the CS is granted in that order.

Coordination and Agreement

➤ Distributed Mutual Exclusion

■ Algorithms for mutual exclusion

✓ The criteria:

- the bandwidth consumed, which is proportional to the number of messages sent in each entry and exit operation;
- the client delay incurred by a process at each entry and exit operation;
- the algorithm's effect upon the throughput of the system.

✓ Some examples of algorithms:

- The central server algorithm
- Ring-Based Algorithm
- Multicast and Logical Clocks
- Maekawa's Voting Algorithm


Coordination and Agreement

➤ Distributed Mutual Exclusion

▪ The central server algorithm

- ✓ Employs a server that grants permission to enter the critical section.

Server managing a mutual exclusion token for a set of processes




Coordination and Agreement

➤ Distributed Mutual Exclusion

■ Ring-Based Algorithm

- ✓ Arrange the processes in a logical ring

A ring of processes transferring a mutual exclusion token




Coordination and Agreement

➤ Distributed Mutual Exclusion

■ Multicast and Logical Clocks

- ✓ processes that require entry to a critical section multicast a request message, and can enter it only when all the other processes have replied to this message

Multicast synchronization



Coordination and Agreement

➤ Election

- An algorithm for choosing a unique process to play a particular role
 - ✓ a process p_i can be
 - a *participant* : is engaged in some run of the election algorithm
 - a *non-participant* : is not currently engaged in any election
 - ✓ Some examples of election algorithms
 - A ring-based election algorithm
 - The bully algorithm

Coordination and Agreement

➤ Election


■ The Bully algorithm

- ✓ Allows processes to crash during an election, although it assumes that message delivery between processes is reliable
- ✓ It assumes that the system is synchronous (uses timeouts to detect a process failure)
- ✓ It assumes that each process knows which processes have higher identifiers, and that it can communicate with all such processes
- ✓ 3 types of message used in this algorithm
 - an *election message* - is sent to announce an election;
 - an *answer message* - is sent in response to an election message
 - a *coordinator message* - is sent to announce the identity of the elected process


Coordination and Agreement

➤ Coordination and Agreement in Group Communication

- *System*: contains a collection of processes, which can communicate reliably over one-to-one channels
- *Processes*: members of groups, may fail only by crashing



Closed group



Open group

Coordination and Agreement

➤ Coordination and Agreement in Group Communication

■ Primitives:

- ✓ ***multicast(g, m)***: sends the message m to all members of group g
- ✓ ***deliver(m)*** : delivers the message m to the calling process
- ✓ ***sender(m)*** : unique identifier of the process that sent the message m
- ✓ ***group(m)***: unique identifier of the group to which the message m was sent

Coordination and Agreement

➤ Coordination and Agreement in Group Communication

■ Basic Multicast

- ✓ Guarantee that a correct process will eventually deliver the message as long as the multicaster does not crash

- ✓ **Primitives:** $B_multicast$, $B_deliver$

- ✓ **Implementation:** Use a reliable one-to-one communication

- ✓ **Unreliable:** Acknowledgments may be dropped

To $B_multicast(g, m)$: for each process $p \in g$, $send(p, m)$;

On $receive(m)$ at p : $B_deliver(m)$ at p .

Coordination and Agreement

➤ Coordination and Agreement in Group Communication

■ Reliable Multicast

■ Properties to satisfy:

- **Integrity:** A correct process P delivers the message m at most once
- **Validity:** If a correct process multicasts a message m , then it will eventually deliver m