

Distributed Transaction Requirements

General characteristics of distributed systems


- Independent failure mode
- No global time
- Inconsistent state

Need to consider:


- how to achieve distributed commitment (or abort)
- how to achieve distributed concurrency control

- A distributed transaction is composed of several sub-transactions, each running on a different site
- Transaction supports ACID properties. Distributed transaction also support ACID properties.
- ~~But some properties~~
- Why do we study Distributed Transactions?
 - Some properties harder to implement
 - Basic single-system techniques not sufficient

Simple Distribute Model



Nested Transaction



If client runs transactions then
Each transaction must complete before proceeding
to next

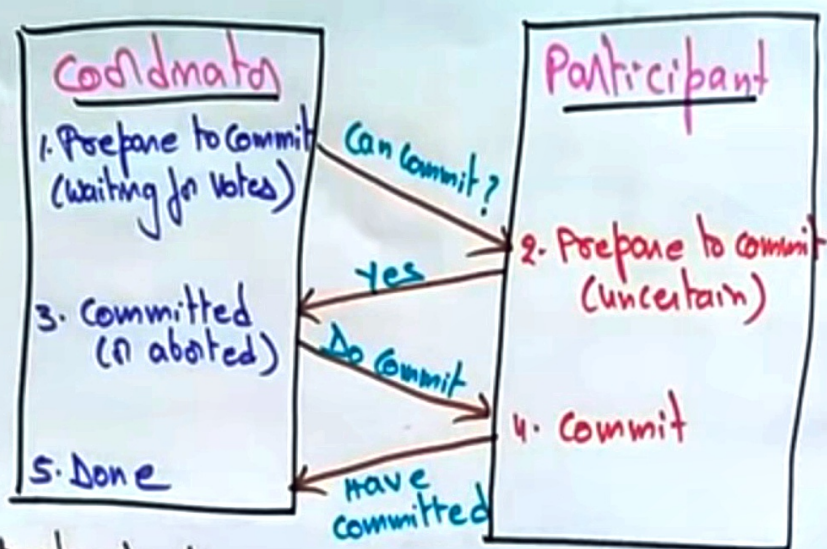
transactions. The coordinator handles all communication with other servers.

Atomic Commit protocols:

- Distribution implies \pm independent failure modes, i.e. one can fail at any time, & others may not discover.
- If one phase commit, client requests commit, but one of the servers may have failed - no way of ensuring durability.
- Instead, commit in 2 phases, thus allowing servers to request abort.

2phase commit

- One coordinator responsible for initiating protocol
- other entities called participants
- If coordinator or participants unable to commit, all participant transaction are aborted.
- Two phases
 - phase 1: The coordinator sends a Can Commit? msg to all participants in transaction
 - phase 2: Implement that decision at all sites.



- Note! - If participant crashes after voted to commit, it can ask coordinator about results of vote.
- Timeouts are used when msg are expected
 - Introduces new state in transaction prepare to commit

