

Machine Learning

Theory & Practice

Module 3: Linear Regression


Lecture 2: Performance Tuning for Linear Regression



Lecture Outline

Topic 1: **R-squared and Adjusted R-squared**

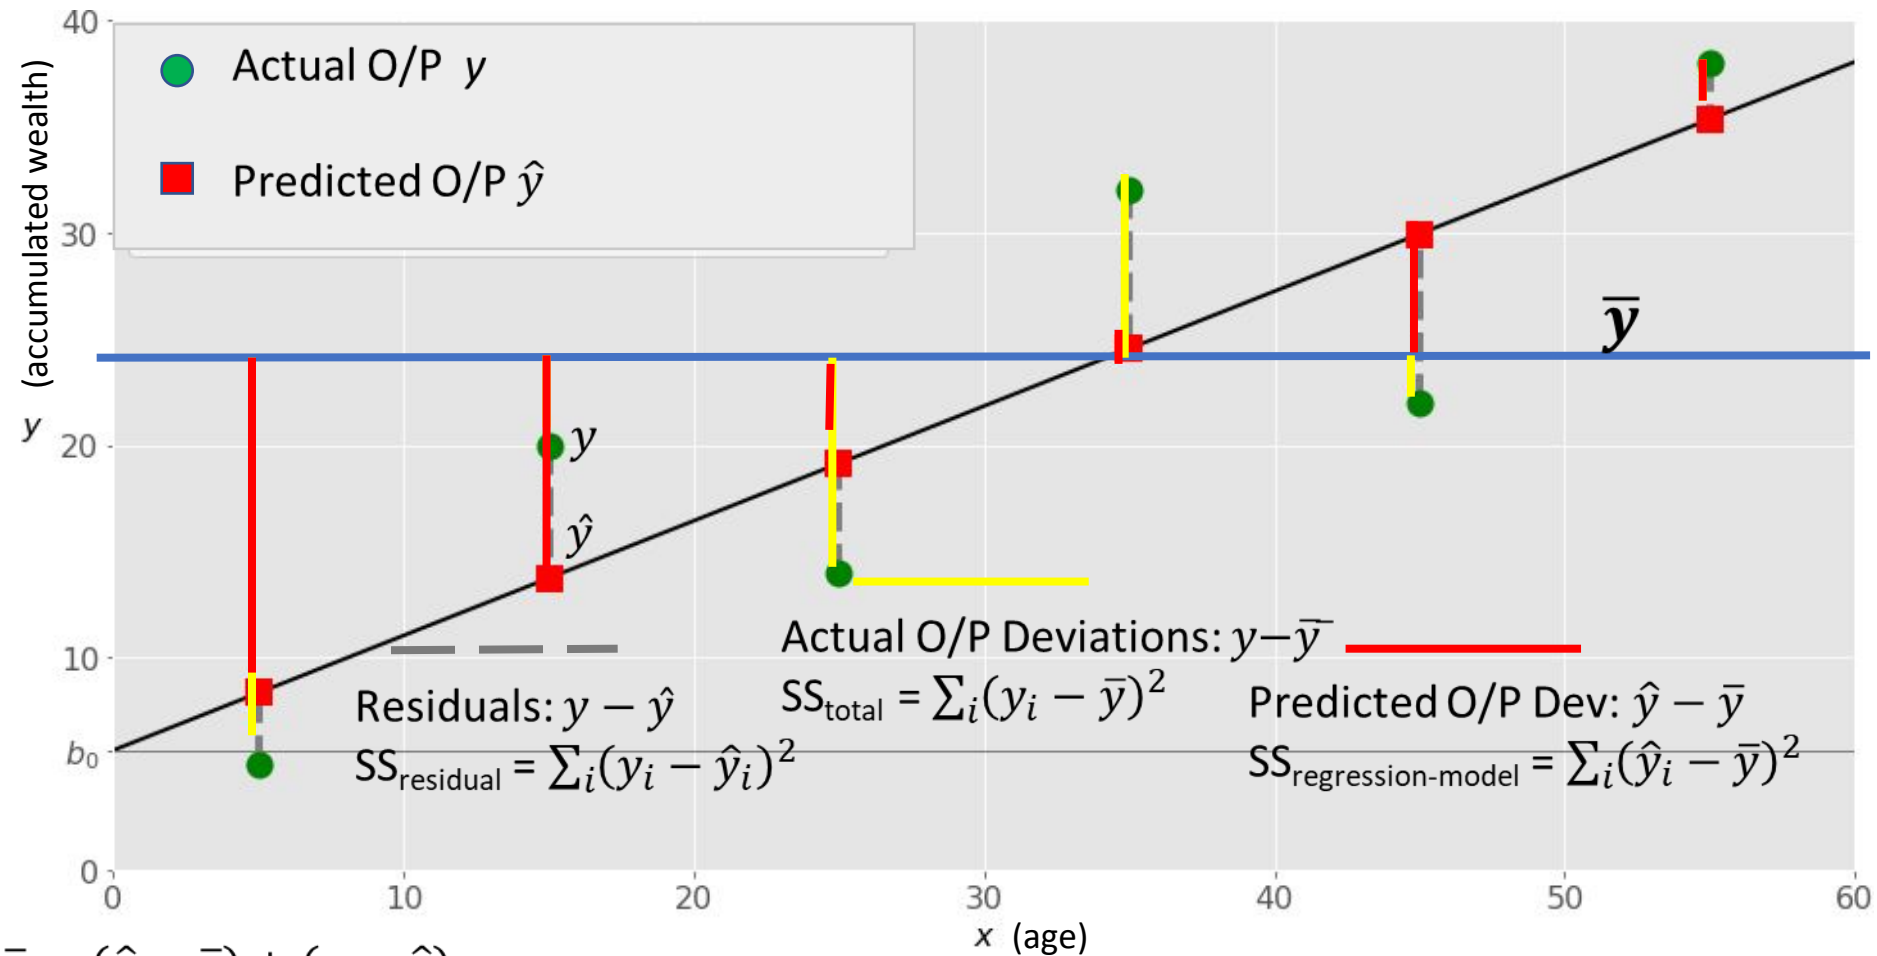
Topic 2: **Regularization**



Topic 1: R-Squared and Adjusted R-Squared:

These are performance metrics that show how much of the output variation is explained by a LR model

Variation in Response



Coefficient of Determination: R-Squared - R^2

Variances:

- Total variance in actual response: $SS_{\text{total}} = \sum_{i=1 \dots n} (y_i - \bar{y})^2$
- Error: $SS_{\text{residual}} = \sum_{i=1 \dots n} (y_i - \hat{y}_i)^2$
- Variance in response given by Model: $SS_{\text{reg-model}} = \sum_{i=1 \dots n} (\hat{y}_i - \bar{y})^2$
- $R^2 = \frac{SS_{\text{reg-model}}}{SS_{\text{total}}}$: proportion of the variation in actual response that is explained by the LR model
- OR, $1 - \frac{SS_{\text{residual}}}{SS_{\text{total}}}$: 1- proportion of unexplained variation (error) in response

Issues with R-squared

-

$$R^2 = 1 - \frac{SS_{residual}}{SS_{total}} = 1 - \frac{\sum_i (y_i - (w_0 + w_1 x_1 + \dots))^2}{\sum_i (y_i - \bar{y})^2}$$

- It is biased: By adding many regressors, and by increasing training time, we get apparently better R^2
- May add insignificant predictors
- May start modelling noise
- Model overfits, and hence will be unable to generalize

Adjusted-R Squared \bar{R}^2

Mean Variances

- $MS_{\text{total}} = \frac{\sum_{i=1 \dots n} (y_i - \bar{y})^2}{n-1}$, n: No. of samples, n-1 : degree of freedom df
- $MS_{\text{regression-model}} = \frac{\sum_{i=1 \dots n} (\hat{y}_i - \bar{y})^2}{p}$, df=p, p= no of regressors (1 for SLR)
- $MS_{\text{residual}} = \frac{\sum_{i=1 \dots n} (y_i - \hat{y}_i)^2}{n-p-1}$, Note: $df_{\text{total}} = df_{\text{regression-model}} + df_{\text{residual}}$
- $\bar{R}^2 = 1 - \frac{MS_{\text{residual}}}{MS_{\text{total}}} = 1 - \frac{SS_{\text{residual}}}{SS_{\text{total}}} \frac{(n-1)}{(n-p-1)}$

Relationship between R^2 & \bar{R}^2

R^2 Versus \bar{R}^2

- LR: $R^2 = 1 - SS_{res} / SS_{tot}$
 - R^2 Always increases as more regressors are added.
 - R^2 is biased estimate
 - R^2 Does not penalize non-significant terms
 - R^2 Is always positive
 - R^2 is not suitable for statistical test of significance of weights
- LR: $\bar{R}^2 = 1 - MS_{res} / MS_{tot}$
 - \bar{R}^2 increases ONLY if an added regressor is significant.
 - \bar{R}^2 is unbiased estimate
 - \bar{R}^2 penalizes non-significant variables
 - \bar{R}^2 can be negative, is always less than R^2
 - \bar{R}^2 is suitable for statistical test of significance of weights



Topic 2: Regularization methods

*REGULARIZATION: PROCESS OF MAKING THE LEARNING
MODEL SIMPLER*

Regularization with Validation


- Make a hierarchy of regressors (X) in terms of their relative importance.
- Add Regressors, one/few at a time, in order of importance.
- Generate an optimized model for each set of regressors using cross validation, and monitor \bar{R}^2
- The point at which \bar{R}^2 reaches a maximum gives the ideal combination of regressors.

Example: Cost of House

Regressors:

1. Floor Area
2. No of bedrooms
3. No of balconies
4. Type of locality
5. Green cover
6. Front door facing which direction
7. Educational background of neighbourhood

$$\bar{R}^2 = 1 - \frac{SS_{residual}}{(n-p-1)} \frac{(n-1)}{SS_{total}}$$




$$\bar{R}^2 = 1 - \frac{SS_{residual}}{SS_{total}} \frac{(n-1)}{(n-p-1)}$$

Regularization with Modified Loss Functions

- Augment Ordinary Least Squares with regularization term:
 - LASSO Regression \Rightarrow L1 Regularization
 - Ridge Regression \Rightarrow L2 Regularization
 - Elastic Net Regularization

Least Absolute Shrinkage & Selection Operator(LASSO): L1 Regularization

Minimize cost function:

1) Ordinary Least Squares

2) Regularization Term

$$\textit{Minimize} \left\{ \sum_{i=1}^M \left(y_i - \sum_{j=0}^p w_j \times x_{ij} \right)^2 + \lambda \sum_{j=0}^p |w_j| \right\}$$

Forcing For some $t > 0$, $\sum_{j=0}^p |w_j| < t$

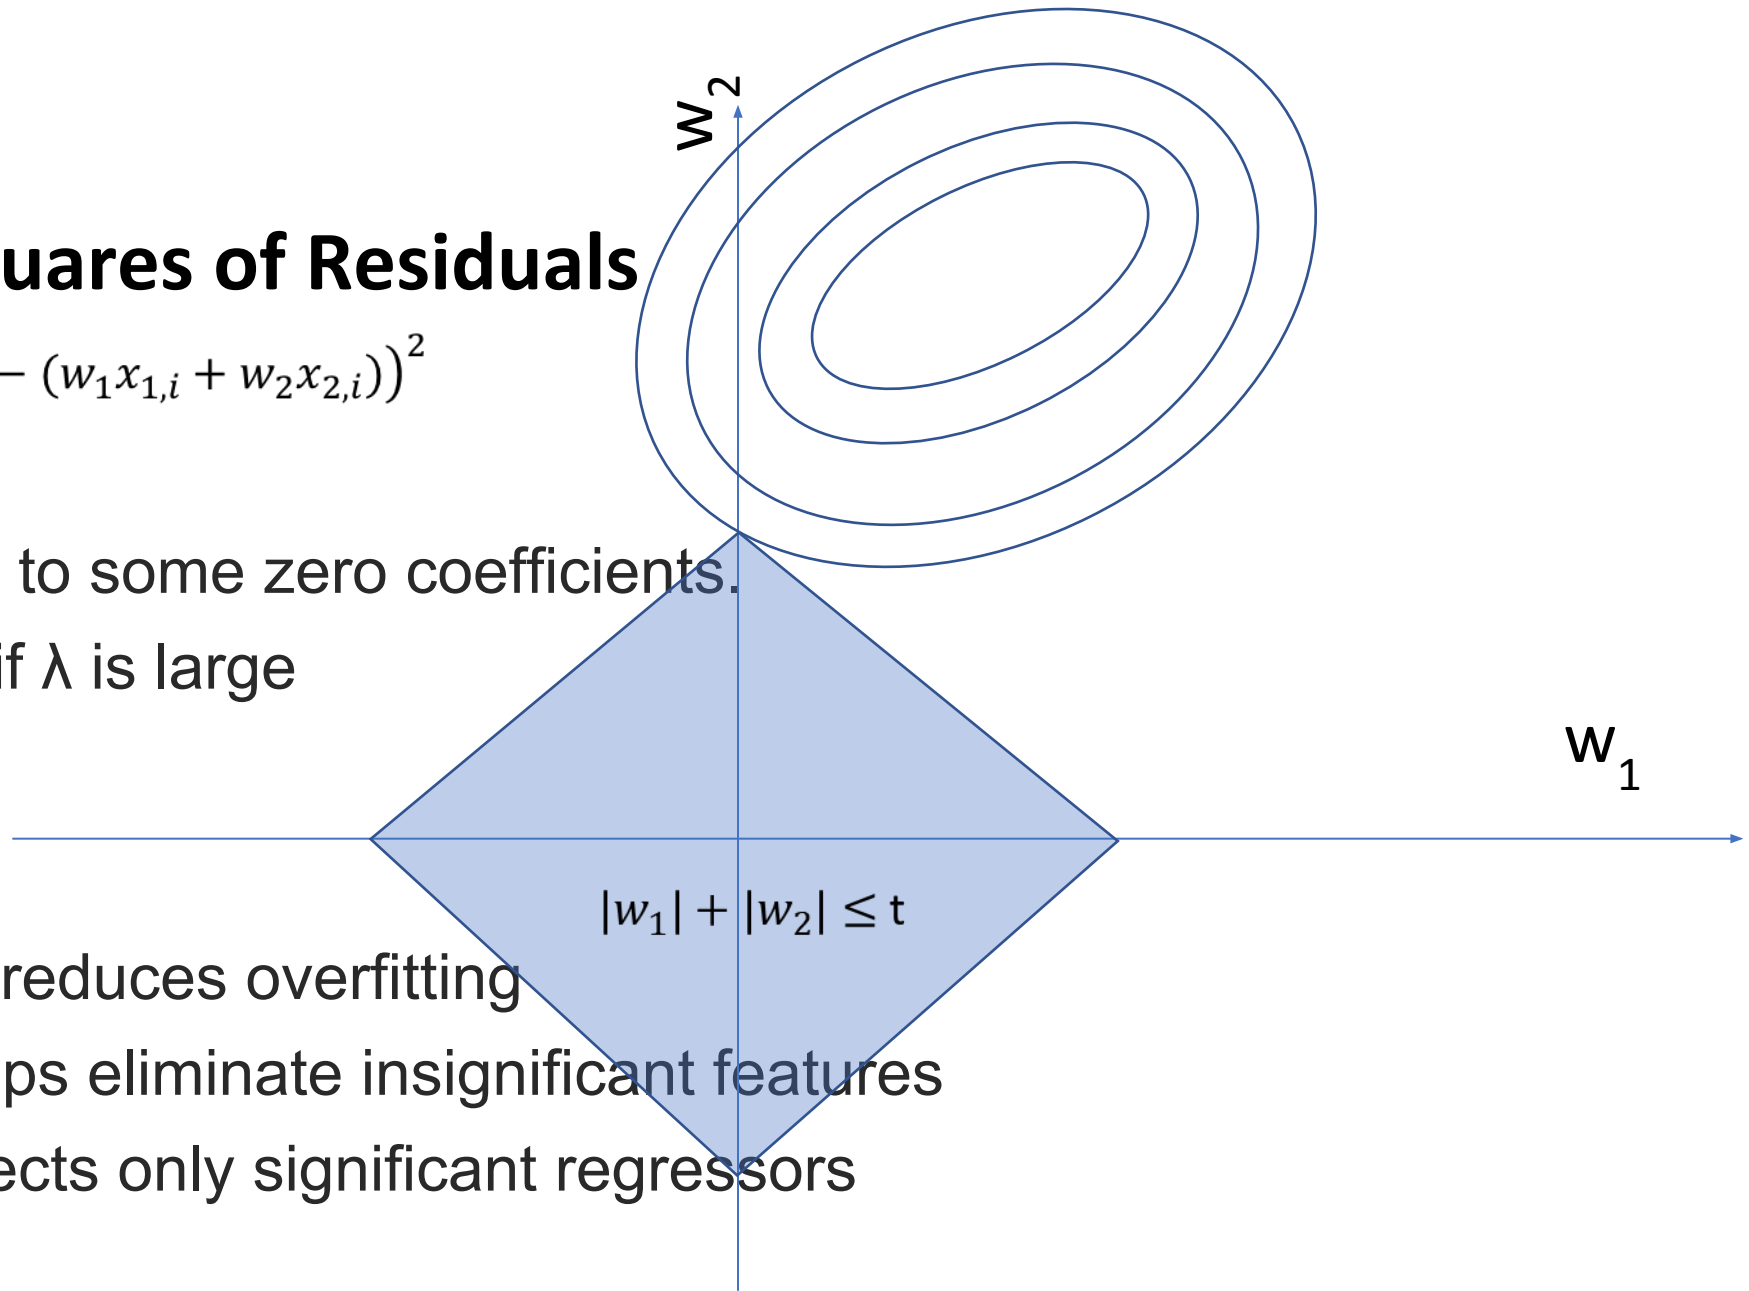
- L1 penalizes regressors by shrinking their weights
- Regressors that contribute little to error reduction are more penalized
- λ is the weighting factor for regularization to tune overfit \leftrightarrow underfit

Sum of squares of Residuals

$$\sum_i (y_i - (w_1 x_{1,i} + w_2 x_{2,i}))^2$$

- L1 can lead to some zero coefficients.
especially if λ is large

- L1 not only reduces overfitting
but also helps eliminate insignificant features
- LASSO selects only significant regressors



Multicollinear Regressors

- The following regressors are highly positively correlated and each apparently impacts property cost.
 - Total Area
 - Number of rooms
 - Free areas: balconies and corridors
 - Size of rooms
- Without regularization, all coefficients would be inflated
- L1 retains only one of them and eliminates the rest

Ridge Regression : L2 Regularization

Minimize cost function: 1) Ordinary Least Squares 2) Regularization Term

$$\textit{Minimize} \left\{ \sum_{i=1}^M \left(y_i - \sum_{j=0}^p w_j \times x_{ij} \right)^2 + \lambda \sum_{j=0}^p w_j^2 \right\}$$

Forcing, For some $c > 0$, $\sum_{j=0}^p w_j^2 < c$


- L2 leads to near zero coefficients

Constraint Region for L2

$$(w_1^2 + w_2^2) \leq t$$

$$SS_{\text{residuals}} = \sum_i (y_i - (w_1 x_{1,i} + w_2 x_{2,i}))^2$$


- L2 handles multiple correlated regressors better



Elastic Net Regularization

- Combines L1 and L2 Regularization
- Each has its own weighting factor
- $OF = Minimize \left\{ \sum_{i=1}^n \left(y_i - \sum_{j=0}^m w_j x_{j,i} \right)^2 + \lambda_1 \sum_{j=0}^m |w_j| + \lambda_2 \sum_{j=0}^m w_j^2 \right\}$
- λ_1 and λ_2 allow:
 - A balance of attribute elimination ability and handling multiple correlated regressors
 - A proper tuning of overfitted model (both 0) to underfitted model (both large)

Constraint Regions in Ridge, LASSO & Elastic Net Regularization



Recap

- Adjusted Coefficient of Determination \bar{R}^2 indicates the proportion of normalized variation in response that is explained in an unbiased manner, by a LR model
- LASSO regression employs L1 regularization by adding “sum of absolute weights” term in Cost function, with weighting factor λ
- If λ is 0, there is no regularization and the model may overfit giving poor generalization
- If λ is large, most weights shrink and the model may underfit

Recap

- Ridge regression employ L2 regularization by adding “sum of squares of weights” to OLS, with weight factor λ
- L2 regularization cannot eliminate any regressor but can appropriately shrink insignificant ones
- Ridge regression handles multiple correlated features by diminishing or enhancing them simultaneously, instead eliminating all but one (as in L1)
- Elastic Net regularization employs a mix of L1 and L2 with their own weighting factors to create a balanced model



Balance is a sense of harmony....

