


Reinforcement Learning

RL is learning from interaction

RL is learning from interaction



S T
A R

Markov Decision Process (MDP)

- Set of states S – actually features vector to describe environment & agent
- Set of actions A that can be taken at each state
- State transition probabilities $p(s' \mid s, a)$. Memory-less property assumed.
- Policy is mapping from States to possible actions $P: S \rightarrow A$. This is the solution of MDP
- Finite MDP if both S and A are finite

Markov Decision Process (MDP)

- Rewards are a set of real numbers or discrete in $S \times A$. Reward Functions:
 - State reward $R(S)$
 - Immediate reward $R(S,a, S')$,
 - $R(S,a)$.
- Discount factor γ in $[0, 1]$ – priority given to future experience. ~ 0.9
- Learning rate: α (~ 0.1)

Lodhi Garden visit – R: State Transition Immediate rewards r_{t+1}

$(S_t \rightarrow S_{t+1})$	S_1 Tomb, hungry, tired	S_2 Garden, hungry, fresh	S_3 Lake, filled, tired	S_4 Eatery, filled, fresh	S_5 Gazebo, hungry, fresh	S_6 Out, filled, tired
S_1 Tomb, hungry, tired	-10	50	70	100	50	0
S_2 Garden, hungry, fresh	20	40	100	100	20	0
S_3 Lake, filled, tired	0	70	20	0	0	50
S_4 Eatery, filled, fresh	100	0	0	0	0	0
S_5 Gazebo, hungry, fresh	0	0	100	100	0	0
S_6 Out, filled, tired	0	0	60	0	0	100

What is the feature vector that defines different states?

Q learning basics

- Q is a quality or utility or Value Function -> helps assess decisions
- Maximizes sum of
 - Immediate reward
 - Projected future reward
- Recursive in nature
- Q must be updated with experience
- Initial Q : all zeros

Initial Utilities – all zeros. Initial State S_3 , Discount factor $\gamma = 0.8$

	S_1 Tomb, hungry, tired	S_2 Garden, hungry, fresh	S_3 Lake, filled, tired	S_4 Eatery, filled, fresh	S_5 Gazebo, hungry, fresh	S_6 Out, filled, tired
S_1 Tomb, hungry, tired	0	0	0	0	0	0
S_2 Garden, hungry, fresh	0	0	0	0	0	0
S_3 Lake, filled, tired	0	0	0	0	0	0
S_4 Eatery, filled, fresh	0	0	0	0	0	0
S_5 Gazebo, hungry, fresh	0	0	0	0	0	0
S_6 Out, filled, tired	0	0	0	0	0	0

Learning by interacting – Episode 1

- Let first action randomly be $P(S_2 | S_1) \rightarrow$ Go to Garden.
- Immediate reward = 70
- $Q \rightarrow$ See next state S_2 = Discount factor $\gamma * \text{Max}\{0,0,0,0,0,0\} = 0$
- Thus new $Q(S_1, S_2) = 70 + 0$, due to only instant reward

$(S_t \rightarrow S_{t+1})$	S_1 Tomb, hungry, tired	S_2 Garden, hungry, fresh	S_3 Lake, filled, tired	S_4 Eatery, filled, fresh	S_5 Gazebo, hungry, fresh	S_6 Out, filled, tired
S_1 Tomb, hungry, tired	-10	50	70	100	50	0
S_2 Garden, hungry, fresh	20	40	100	100	20	0
S_3 Lake, filled, tired	0	70	20	0	0	50
S_4 Eatery, filled, fresh	100	0	0	0	0	0
S_5 Gazebo, hungry, fresh	0	0	100	100	0	0
S_6 Out, filled, tired	0	0	60	0	0	100

Updated Utility: Q'

	S ₁ Tomb, hungry, tired	S ₂ Garden, hungry, fresh	S ₃ Lake, filled, tired	S ₄ Eatery, filled, fresh	S ₅ Gazebo, hungry, fresh	S ₆ Out, filled, tired
S ₁ Tomb, hungry, tired	0	0	0	0	0	0
S ₂ Garden, hungry, fresh	0	0	0	0	0	0
S ₃ Lake, filled, tired	0	70	0	0	0	0
S ₄ Eatery, filled, fresh	0	0	0	0	0	0
S ₅ Gazebo, hungry, fresh	0	0	0	0	0	0
S ₆ Out, filled, tired	0	0	0	0	0	0

Episode 2: Initial state = S_6 Outside

- Exploration: Randomly move to S_3

Reward of 6th row

S_6 Out, filled, tired	0	0	60	0	0	100
--------------------------	---	---	----	---	---	-----

- Reward = 60.

Q of 3rd row

S_3 Lake, filled, tired	0	70	0	0	0	0
---------------------------	---	----	---	---	---	---

- $Q = 0.8 * \max\{70, 0\} = 56$
- Updated $Q(S_6, S_3) = 60 + 56 = 116$
- Repeat episodes to reach convergence

Updated Utility

	S ₁ Tomb, hungry, tired	S ₂ Garden, hungry, fresh	S ₃ Lake, filled, tired	S ₄ Eatery, filled, fresh	S ₅ Gazebo, hungry, fresh	S ₆ Out, filled, tired
S ₁ Tomb, hungry, tired	0	0	0	0	0	0
S ₂ Garden, hungry, fresh	0	0	0	0	0	0
S₃ Lake, filled, tired	0	70	0	0	0	0
S ₄ Eatery, filled, fresh	0	0	0	0	0	0
S ₅ Gazebo, hungry, fresh	0	0	0	0	0	0
S ₆ Out, filled, tired	0	0	116	0	0	0

Q learning algorithm

For each (s,a) , initialize $\hat{Q}(s, a) = 0$

Observe Current State

Do Forever:

1. Select action a and execute it
2. Receive Immediate Reward r
3. Observe new state s'
4. Update Utility Table for $\hat{Q}(s, a)$ as:

$$\hat{Q}(s, a) = r + \gamma \times \text{Max}_{a'} Q(s', a')$$

5. Enter New State

Learning Rate

$$\begin{aligned} Q[\text{state}, \text{action}] = & \\ & Q[\text{state}, \text{action}] + \\ & \text{learning Rate} * \\ & (\text{reward} + \text{Discount Factor} * \max_i \{Q[\text{new_state}, S_i]\} \\ & - Q[\text{state}, \text{action}]) \end{aligned}$$

Discount Factor γ weights future rewards possible, farther the future event, lesser the reward.

Reinforcement Learning:

Each State is a combination of features

Can features change to produce a new state?

- Can an ML agent change features?

=> Yes, provided it interacts with the environment that produces data

- Can features change every now and then by itself?

=> When many entities participate and interact with environment (AI agents/Humans/Nature)

- What applications need this scenario?

=> Where action is needed.

=> Where environment / features are dynamically changing

=> Where there is no training data or human expertise

Applications of RL

- Autonomous car.....What features does it change?
- Trading does time series or regression analysis do this? What features change?
- Making decisions at real time using multimedia data A swarm of robots operating
 - In manufacturing
 - In dangerous and unknown areas
- NLP – Text summarization, machine translation, Question Answering, chat-box, dialogue generation
 - supervised DL models to predict words
 - RL through rewards when to look for more words / where to look for important words
 - Need to define reward in terms of linguistic quality parameters such as cohesiveness, understandability etc.
- Dynamic Healthcare – diagnosis / treatment / drugs manufacture/ medical policies
- Auction bidding

What is Reinforcement Learning?

- Agent interacts with Environment
- RL ? Learning from interaction with an environment
- Environment state ? feature-vector
- Long term goal ? maximize predicted cumulative future rewards
- Agent must be able to partially/fully sense the environment state
- Take actions ? change environment state
- Tradeoff between exploration and exploitation
- Needs lots of training

Deep Reinforcement

Learning?

- A deep neural network is used to develop either a policy or a value function (state to action/ Q-value) or learn features in complex scenarios
- Deep neural networks require lots of real/simulated interaction with the environment to learn
- Lots of trials/interactions is possible in simulated environments
- We can easily parallelise the trials/interaction in simulated environments

Model-free versus

Model-based

- A model of the environment allows inferences to be made about how the environment will behave
- Example: Given a state and an action to be taken while in that state, the model could ***predict the next state and the next reward***
- Models are used for planning, which means deciding on a course of action by considering possible future situations before they are experienced
- Model-based methods use models and planning. Think of this as modelling the dynamics $p(s' \mid s, a)$
- Model-free methods learn exclusively from trial-and-error (i.e. no modelling of the environment)